

Artificial Neural Networks

Inteligencia Artificial en los Sistemas de Control Autónomo
Máster Universitario en Ingeniería Industrial

Departamento de Automática

Objectives

1. Describe biological neurons and networks
2. Basics of artificial neurons and networks
3. Understand the role of training in ANNs
4. Strengths and weaknesses of ANNs

Bibliography


- A. Tettamanzi, M. Tomassini. Soft Computing. Integrating Evolutionary, Neural, and Fuzzy Systems. Springer-Verlag. 2001
- McCulloch, W. and Pitts, W. (1943). A logical calculus of the ideas immanent in nervous activity. Bulletin of Mathematical Biophysics, 7:115 - 133.
- Rosenblatt, Frank. (1958). The Perceptron: A Probabilistic Model for Information Storage and Organization in the Brain. Psychological Review, 65:386-408

Table of Contents

Introduction

History

- 1888 Ramón y Cajal. Discovery of biological neurons
- 1943 McCulloch & Pitts. First neural network designers
- 1949 Hebb. First learning rule
- 1958 Rosenblatt. Perceptron
- 1969 Minsky & Papert. Perceptron limitation - Death of ANN
- 1986 Rumelhart et al. Re-emergence of ANN: Backpropagation
- 201X Convolutional Neural Networks (CNNs) - Deep learning
- 2014 Goodfellow et al. Generative Adversarial Networks (GANs)



Introduction

Structure of neurons (I)


ANIMAL	NEURONS
Sponge	0
Roundworm	302
Jellyfish	800
Ant	250,000
Cockroach	1,000,000
Frog	16,000,000
Mouse	71,000,000
Cat	760,000,000
Macaque	6,376,000,000
Human	86,000,000,000
Elephant	267,000,000,000

Human brain

Neuron switching time: 0.001 s

Synapsis: 10-100 thousand

Scene recognition time: 0.1 s



(Source)



Introduction

Structure of neurons (II)

A neuron has a cell body ...

- ... a branching input structure (dendrite) and
- ... a branching output structure (axon)

Axons connect to dendrites via **synapses**



Introduction

Structure of neurons (III)

A neuron only fires if its input signal exceeds a threshold


- Good connections allowing a large signal
- Slight connections allowing a weak signal
- Synapses may be either excitatory or inhibitory

Synapses vary in strength

- Biological learning involves setting that strength

Artificial neurons

Definition (I)



a_i Normalized input ($0 \leq a_i \leq 1$)

w_i Weight of input j

θ Threshold

g Activation function

Neuron model
(perceptron)

$$\gamma = g \left(\sum_{i=1}^n w_i a_i \right)$$

Artificial neurons

Definition (II)

- Each neuron has a threshold value
- Each neuron has weighted inputs
- The input signals form a weighted sum
- If the activation level exceeds the threshold, the neuron activates


Artificial neurons

Definition (III)

The idealized activation function is a step function


$$g(x) = \begin{cases} 1 & \text{if } \sum_{i=1}^N w_i x_i > \theta \\ 0 & \text{otherwise} \end{cases}$$

The step function is rarely used in practice




Artificial neurons


Logical gates with a neuron



a_1	a_2	γ
0	0	0
0	1	0
1	0	0
1	1	1



a_1	a_2	γ
0	0	0
0	1	1
1	0	0
1	1	1




a_1	γ
0	1
1	0

(A neuron in Excel)

Artificial neurons

Definition of neuron (alternative version)



a_i Normalized input ($0 \leq a_i \leq 1$)

w_i Weight of input j

w_0 Bias

g Activation function


Neuron model

$$\gamma = g \left(\sum_{i=0}^n w_i a_i \right)$$

Artificial neurons

Example of biased neuron

AND logical gate with a biased input




a_0	a_1	a_2	Output
I	O	O	O
I	O	I	O
I	I	O	O
I	I	I	I

Artificial neurons

Activation functions: Sigmoid function

- S-shaped, continuous and everywhere differentiable
- Asymptotically approach saturation points
- Derivative fast computation
- Range $\in [0, 1]$



Sigmoid function


$$g(x) = \frac{1}{1 + e^{-x}}$$

$$g'(x) = g(x)(1 - g(x))$$

Artificial neurons

Activation functions: Tanh function

- Asymptotically approach saturation points
- Range $\in [-1, 1]$
- Bigger derivative than sigmoid (faster training)



Tanh function

$$g(x) = \tanh(x) = \frac{2}{1 + e^{-2x}} - 1$$

$$g'(x) = 1 - g(x)^2$$

Artificial neurons

Activation functions: Softmax function

- Generalization of the logistic function
- Usually used in the output layer in classification problems
- Asymptotically approach saturation points


Softmax function

$$g(\mathbf{z})_j = \frac{e^{z_j}}{\sum_{k=1}^K e^{z_k}} \text{ for } j = 1, \dots, K$$


with \mathbf{z} a K -dimensional vector

Artificial neurons


Other activation functions




(a) Linear function




(b) Step function




(c) Ramp function



(d) Sigmoid function



(e) Hyperbolic tangent function




(f) Gaussian function

(Source)

Artificial neurons

Learning limits (I)




Problem: A single neuron only can solve linearly separable problems

Artificial neurons

Learning limits (II)

XOR cannot be implemented with a neuron



Solution: Neuronal networks

Artificial Neural Networks

Definition (I)

- A very much simplified version of biological nerve systems
- A set of nodes (neurons)
 - Each node has input and output
 - Each node performs a simple computation
- Weighted connections between nodes
 - Connectivity gives the structure of the net
 - What can be computed by an ANN is primarily determined by the connections and their weights
- It can recognize patterns, learn and generalize

Artificial Neural Networks

Definition (II)

ANN properties

- Noise tolerance
- General function approximator

Machine Learning tasks

- Supervised learning (classification and regression)
- Unsupervised learning (known as **self-organizing maps** in ANN terminology)

Many topologies

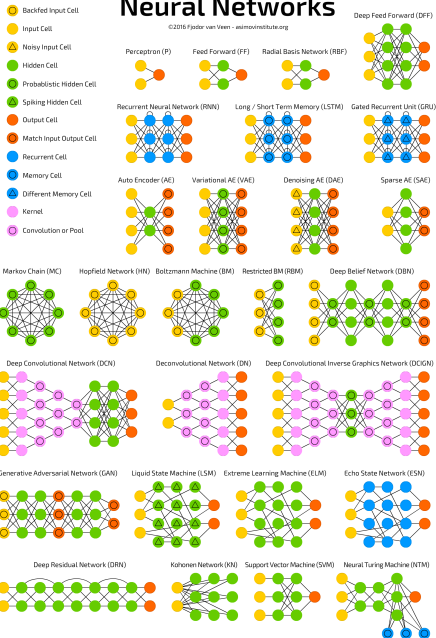
- Acyclic, recurrent (cyclic), modular, etc
- Feed Forward networks (MLPs)

Human readability less important than performance

A mostly complete chart of

Neural Networks


©2016 P.Jodier van Veen - aulmoulinstitute.org



(More info)

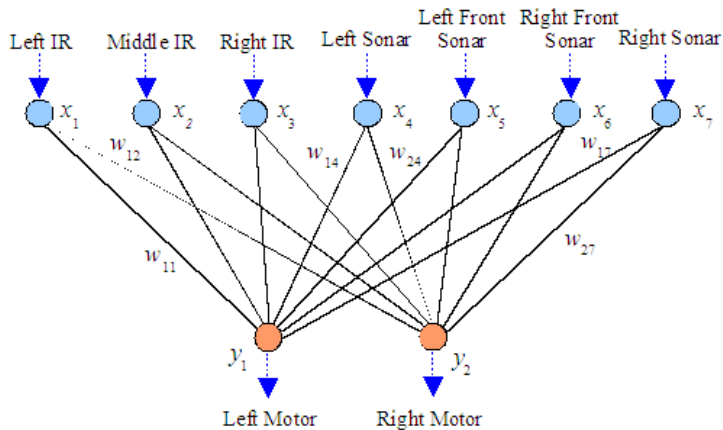
Artificial Neural Networks

Application examples (I)



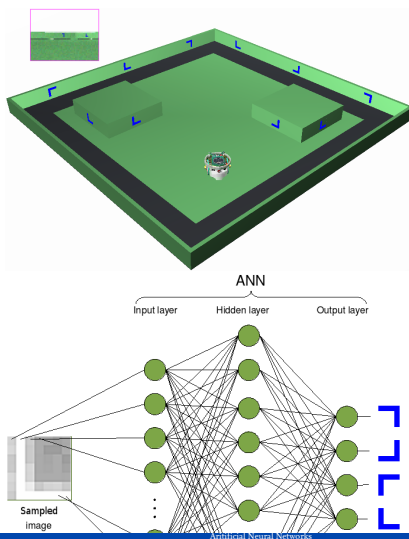
Artificial Neural Networks

Application examples (II)





Artificial Neural Networks

Application examples (III)



Artificial Neural Networks

Application examples (III)



(Source)

Feedforward networks

Definition (I)


Neurons are arranged in **layers**

Input Which consists of any normalized data

Output Which are the net outcome

Hidden (Optional) No direct interaction

Also known as **multilayer perceptron (MLP)**



Feedforward networks

Definition (II)

The input layer

- Introduces input values into the network
- No activation function or other processing

The hidden layer(s)


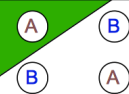
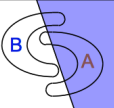

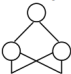
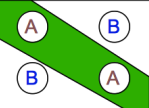
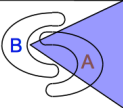
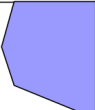
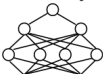
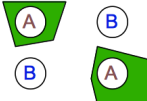

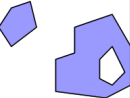
- Perform classification of features
- Two hidden layers are sufficient to solve any problem
- Features imply more layers may be better

The output layer

- Functionally just like the hidden layers
- Outputs are passed on to the world outside the neural network

Feedforward Networks


Separability

Structure	Types of Decision Regions	Exclusive-OR Problem	Classes with Meshed regions	Most General Region Shapes
Single-Layer 	Half Plane Bounded By Hyperplane			
Two-Layer 	Convex Open Or Closed Regions			
Three-Layer 	Arbitrary (Complexity Limited by No. of Nodes)			

(Online demo)

Feedforward Networks

Bias in a MLP



(Source)


Training algorithms

Problem statement (I)

ANN can perform different tasks

- Classification, regression, others

Classification (or supervised learning) uses a training set



A ₁	A ₂	A ₃	O	Y
1,1	2,5	4,5	0,2	-0,1
0,9	2,4	1,2	0,5	0,4
1,0	2,0	9,9	0,4	1,2

Toss function: Measure of the error

- Usually mean squared error (mse): $E = \frac{1}{2}(\gamma - o)^2 = f(w)$
- Y and O are the desired and observed outputs

Training algorithms

Problem statement (II)

$$E = \frac{1}{2} \text{Err}^2 = \frac{1}{2} \left[\gamma - g \left(\sum_{j=0}^n w_j x_j \right) \right]^2$$

where

γ Desired output

w_j Weight connection j

x_j Input j

Problem: Determine w that minimize $f(w)$

- This is a classical optimization problem
- Any optimization algorithm can be used
- ... in AI, optimization means search

Training algorithms

Gradient Descent Algorithm (I)

Given the error

$$E = \frac{1}{2} \text{Err}^2$$

Take partial derivatives

$$\begin{aligned} \frac{\partial E}{\partial w_j} &= \text{Err} \frac{\partial \text{Err}}{\partial w_j} \\ &= \text{Err} \frac{\partial}{\partial w_j} g \left(y - \sum_{j=0}^n w_j x_j \right) \\ &= -\text{Err} \times g'(w) \times x_j \end{aligned}$$

Training algorithms

Gradient Descent Algorithm (II)

Weight update

$$w_j^{k+1} = w_j^k + \alpha \times \text{Err} \times g'(w) \times x_j$$

with

α Learning rate ($|\alpha| < 1$)

err Difference desired and current output

g' Derivate of activation function

x_j Input j


Each iteration is named **epoch**

Learning algorithm (single neuron)

1. Apply input signal and compute outout
2. If output == desired output, do nothing
3. If output < desired output, increase weights
4. If output > desired output, decrease weights



Training algorithms

Gradient Descent Algorithm (III)



(Source)

Learning rate



(Source)

Training algorithms

Stochastic Gradient Descent (I)

SDG approximates the gradient sampling the dataset

On-line One sample

Mini-batch Several samples

Batch All the samples (Gradient Descent)


Weights update rule: $w^{k+1} = w^k - \alpha \nabla g(\text{in})$

- where α is the learning rate

SGD is slow and prone to local minima

Training algorithms

Stochastic Gradient Descent (II)



Usually, a **momentum** is introduced: $w^{k+1} = w^k - \alpha z^{k+1}$, where $z^{k+1} = \beta z^k + \nabla g(\text{in})$

- α is the learning rate
- β is the momentum strength
- If $\beta = 0$ then gradient descent


(On-line demo)

Training algorithms

Other optimization algorithms

Other optimization algorithms

- Newton's method
- Quasi-Newton's method
- Levenberg-Marquardt method
- Conjugate Gradient



Training algorithms


Backpropagation algorithm (I)

Efficient learning algorithm for multilayer perceptrons. Three steps

1. **Feed-forward step.** Feed input, compute output and error
2. **Feed-backward step.** Compute individual contribution to error
3. **Adjust weights.** Modify weights to minimize error: Input, output and hidden layers

Training algorithms

Backpropagation algorithm (II)



Output layer: Same as single neuron

$$W_j \leftarrow W_j + \alpha \times \text{Err} \times g'(\text{in}) \times x_j$$

Define modified error as
 $\Delta_i = \text{Err}_i \times g'(\text{in}_i)$, then

$$W_{j,i} \leftarrow W_{j,i} + \alpha \times a_j \times \Delta_i$$

Training algorithms

Backpropagation algorithm (III)

Hidden layer: Propagate error

$$W_{k,j} \leftarrow W_{k,j} + \alpha \times a_k \times \Delta_j$$

where

$$\Delta_j = g'(\text{in}) \times \sum_i W_{j,i} \Delta_i$$

Backpropagation algorithm


1. Compute output
2. Compute output error Δ
3. For each layer, repeat the following steps
 - 3.1 Propagate Delta backwards
 - 3.2 Update weights between two layers

Training algorithms

Learning problems

Potential problems

- Local minima
- Flat plateau
- Oscillation
- Missing good minima



Training algorithms

Learning problems: Under and overfitting (I)

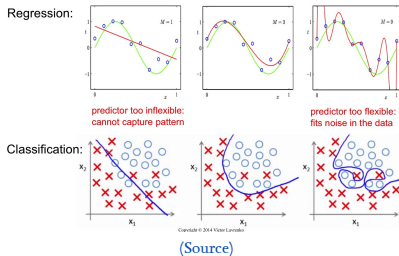
Underfitting: Does not learn

- Topology too simple

Overfitting: Memorizes samples

- Topology too complex
- Perhaps, the most serious concern in ML
- The net fails when exposed to new data

Under- and Over-fitting examples




Training algorithms

Learning problems: Under and overfitting (II)

Solution: Evaluate generalization capabilities

- Split training and validation sets and measure errors



(Source)

Acknowledgements

- Jesus Aguilar Ruiz, Pablo de Olavide, Seville, Spain
- Daniel Rodríguez, Universidad de Alcalá, Alcalá de Henares, Spain