

# Artificial Neural Networks

Inteligencia Artificial en los Sistemas de Control Autónomo  
Máster en Ciencia y Tecnología desde el Espacio

Departamento de Automática

## Objectives

1. Describe biological neurons and networks
2. Basics of artificial neurons and networks
3. Understand the role of training in ANNs
4. Strengths and weaknesses of ANNs

## Bibliography

- A. Tettamanzi, M. Tomassini. Soft Computing. Integrating Evolutionary, Neural, and Fuzzy Systems. Springer-Verlag. 2001
- McCulloch, W. and Pitts, W. (1943). A logical calculus of the ideas immanent in nervous activity. Bulletin of Mathematical Biophysics, 7:115 - 133.
- Rosenblatt, Frank. (1958). The Perceptron: A Probabilistic Model for Information Storage and Organization in the Brain. Psychological Review, 65:386-408

# Table of Contents

## 1. Introduction

- Machine Learning tasks
- History
- Structure of neurons

## 2. Artificial neurons

- Definition
- Logical gates with a neuron
- Activation functions
- Learning limits

## 3. Artificial Neural Networks

- Definition
- Application examples
- Separability
- Topologies
- Layered networks
- Feedforward networks

## 4. Training algorithms

- Problem statement
- Gradient descent algorithm
- Stochastic Gradient Descent Algorithm
- Other optimization algorithms
- Backpropagation algorithm
- Learning problems

# Introduction

## Machine Learning (I)

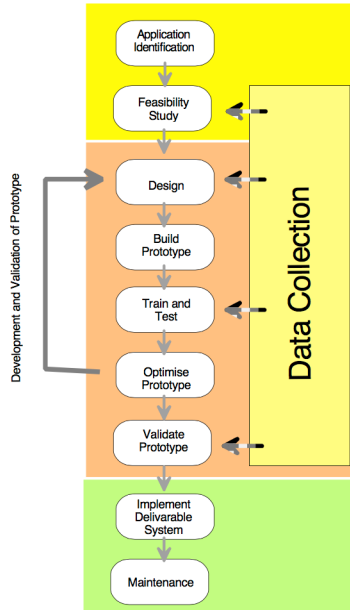
Machine Learning studies how to build data-driven models

- **Supervised learning** The dataset contains output examples
  - **Classification** The output is a category
  - **Regression** The output is a number
- **Unsupervised learning** (or clustering) The dataset does not contain output examples
- **Reinforcement learning** Maximize reward

ML is data-driven

# Introduction

## Machine Learning (II)



# Introduction

## History

1943-McCulloch & Pitts First neural network designers

1949-Hebb First learning rule

1958-Rosenblatt Perceptron

1969-Minsky & Papert Perceptron limitation - Death of ANN

1986 - Rumelhart et al. Re-emergence of ANN: Backpropagation

2012 - Krizhevsky Convolutional Neural Networks - Deep learning

## Introduction

## Structure of neurons (I)


ANIMAL	NEURONS
Sponge	0
Roundworm	302
Jellyfish	800
Ant	250,000
Cockroach	1,000,000
Frog	16,000,000
Mouse	71,000,000
Cat	760,000,000
Macaque	6,376,000,000
Human	86,000,000,000
Elephant	267,000,000,000

## Human brain

Neuron switching time: 0.001 s

Synapsis: 10-100 thousand

Scene recognition time: 0.1 s



(Source)


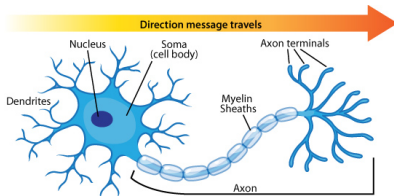
# Introduction

## Structure of neurons (II)

A neuron has a cell body ...

- ... a branching input structure (dendrite) and
- ... a branching output structure (axon)

Axons connect to dendrites via synapses





# Introduction

## Structure of neurons (III)


A neuron only fires if its input signal exceeds a threshold

- Good connections allowing a large signal
- Slight connections allowing a weak signal
- Synapses may be either excitatory or inhibitory

Synapses vary in strength

# Artificial neurons

## Definition (I)



$a_i$  Normalized input ( $0 \leq a_i \leq 1$ )

$w_i$  Weight of input  $j$  ( $0 \leq w_i \leq 1$ )

$\theta$  Threshold

$g$  Activation function

### Neuron model

$$\gamma = g \left( \sum_{i=1}^n w_i a_i \right)$$

# Artificial neurons

## Definition (II)

- Each neuron has a threshold value
- Each neuron has weighted inputs
- The input signals form a weighted sum
- If the activation level exceeds the threshold, the neuron activates


# Artificial neurons

## Definition (III)

The idealized activation function is a step function

$$g(x) = \begin{cases} 1 & \text{if } \sum_{i=1}^N w_i x_i > \theta \\ 0 & \text{otherwise} \end{cases}$$

The step function is rarely used in practice




# Artificial neurons

## Logical gates with a neuron


A neuron can implement a logical gate

AND




$a_1$	$a_2$	$\gamma$
0	0	0
0	1	0
1	0	0
1	1	1

OR



$a_1$	$a_2$	$\gamma$
0	0	0
0	1	1
1	0	0
1	1	1


NOT



$a_1$	$\gamma$
0	1
1	0

# Artificial neurons

## Definition of neuron (alternative version)



$a_i$  Normalized input ( $0 \leq a_i \leq 1$ )

$w_i$  Weight of input  $j$  ( $0 \leq w_i \leq 1$ )

$w_0$  Bias

$g$  Activation function


### Neuron model

$$\gamma = g \left( \sum_{i=0}^n w_i a_i \right)$$

# Artificial neurons

## Example of biased neuron


AND logical gate with a biased input



$a_0$	$a_1$	$a_2$	Summation	Output
-1	o	o	$(-1 * 0,3) + (0 * 0,5) + (0 * 0,4) = -0,3$	o
-1	o	1	$(-1 * 0,3) + (0 * 0,5) + (1 * 0,4) = -0,7$	o
-1	1	o	$(-1 * 0,3) + (1 * 0,5) + (0 * 0,4) = 0,2$	1
-1	1	1	$(-1 * 0,3) + (1 * 0,5) + (1 * 0,4) = -0,2$	o

# Artificial neurons

## Activation functions




(Source)



# Artificial neurons

## Activation functions: Sigmoid function

- S-shaped, continuous and everywhere differentiable
- Asymptotically approach saturation points
- Derivative fast computation
- Range  $\in [0, 1]$



### Sigmoid function


$$g(x) = \frac{1}{1 + e^{-x}}$$

$$g'(x) = g(x)(1 - g(x))$$

# Artificial neurons

## Activation functions: Tanh function

- Asymptotically approach saturation points
- Range  $\in [-1, 1]$
- Bigger derivative than sigmoid (faster training)



### Tanh function

$$g(x) = \tanh(x) = \frac{2}{1 + e^{-2x}} - 1$$

$$g'(x) = 1 - g(x)^2$$

# Artificial neurons

## Activation functions: Softmax function

- Generalization of the logistic function
- Usually used in the output layer in classification problems
- Asymptotically approach saturation points


### Softmax function

$$g(\mathbf{z})_j = \frac{e^{z_j}}{\sum_{k=1}^K e^{z_k}} \text{ for } j = 1, \dots, K$$

with  $\mathbf{z}$  a K-dimensional vector

# Artificial neurons

## Learning limits (I)




Problem: A single neuron only can solve linearly separable problems

# Artificial neurons

## Learning limits (II)

XOR cannot be implemented with a neuron



Solution: Neuronal networks

# Artificial Neural Networks

## Definition (I)

- A very much simplified version of biological nerve systems
- A set of nodes (neurons)
  - Each node has input and output
  - Each node performs a simple computation
- Weighted connections between nodes
  - Connectivity gives the structure of the net
  - What can be computed by an ANN is primarily determined by the connections and their weights
- It can recognize patterns, learn and generalize

# Artificial Neural Networks

## Definition (II)

### ANN properties

- Noise tolerance
- General function approximator

### Machine Learning tasks

- Supervised learning (classification and regression)
- Unsupervised learning (known as **self-organizing maps** in ANN terminology)
  - Autoencoders

### Application examples:

- Robotics, vehicle control, computer vision, videogames, spam filtering

Human readability less important than performance

# Artificial Neural Networks


## Definition (III)

In order to learn, it needs at least two components

Inputs Which consists of any normalized information

Outputs Which are the outcome arrived

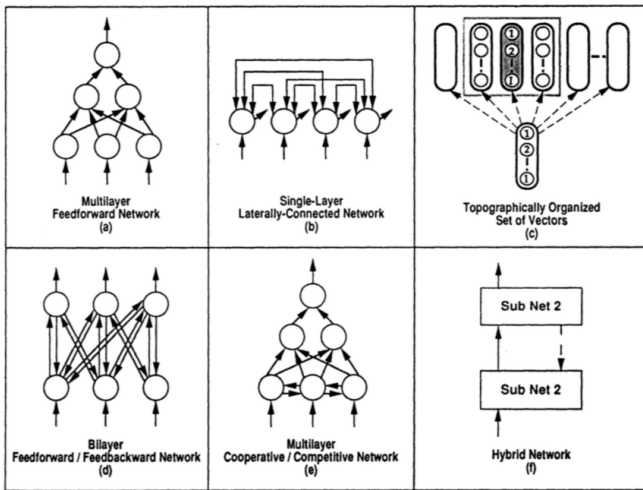
Hidden nodes (Optional) No direct interaction





# Artificial Neural Networks

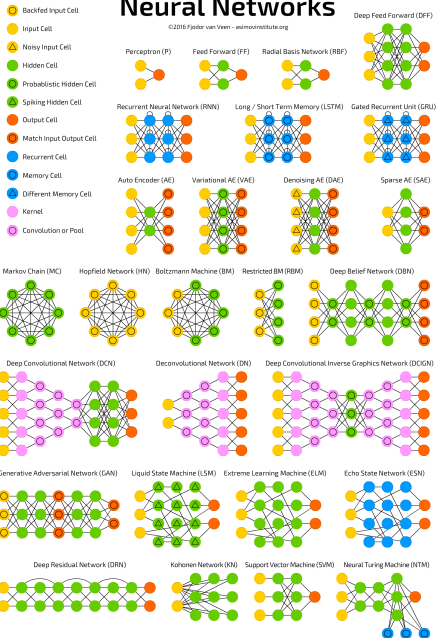
## Definition (IV)



A mostly complete chart of

# Neural Networks


©2016 P.Jodan van Veen - aalmovimstitute.org



(More info)

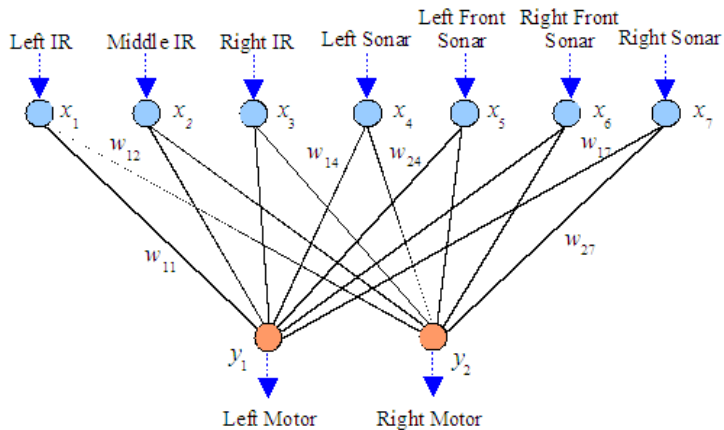
# Artificial Neural Networks

## Application examples (I)




# Artificial Neural Networks

## Application examples (II)

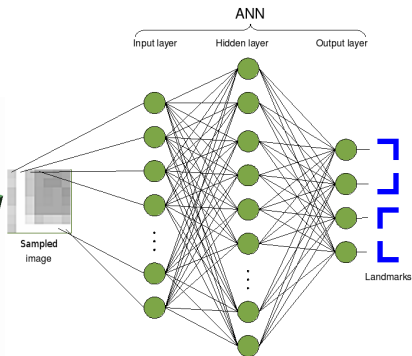


# Artificial Neural Networks

## Application examples (III)


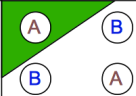
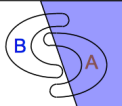

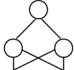
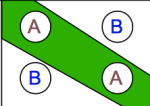
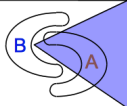
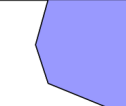
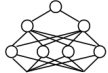
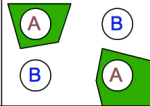

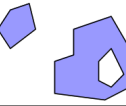


(Source)



# Artificial Neural Networks

## Separability

Structure	Types of Decision Regions	Exclusive-OR Problem	Classes with Meshed regions	Most General Region Shapes
<b>Single-Layer</b> 	<b>Half Plane Bounded By Hyperplane</b>			
<b>Two-Layer</b> 	<b>Convex Open Or Closed Regions</b>			
<b>Three-Layer</b> 	<b>Arbitrary (Complexity Limited by No. of Nodes)</b>			

(Demo online)

# Artificial Neural Networks

## Topologies (I)

### Acyclic Networks

- Without directed cycles
- Easy to analyze

### Recurrent Networks

- With directed cycles
- Much harder to analyze
- Potentially unstable

### Modular nets


- Consists of several modules
- Each module is itself an ANN
- Sparse connections between modules

# Artificial Neural Networks

## Topologies (II)

### Asymmetric fully connected networks

- Every node is connected to every other node
- Connection may be excitatory (positive), inhibitory (negative), or irrelevant (o)
- Most general



### Symmetric fully connected nets


- Weights are symmetric ( $w_{ij} = w_{ji}$ )



# Network architecture

## Layered networks (I)

- Nodes are partitioned into subsets, called layers
- No connections from nodes in layer  $j$  to those in layer  $k$  if  $j > k$





- Inputs are applied to nodes in layer 0
- Nodes in input layer without computation

# Network architecture

## Layered networks (II)

**Perceptron:** ANN whose input is directly connected with its output



# Network architecture

## Layered networks (III)

### The input layer

- Introduces input values into the network
- No activation function or other processing

### The hidden layer(s)

- Perform classification of features
- Two hidden layers are sufficient to solve any problem
- Features imply more layers may be better


### The output layer

- Functionally just like the hidden layers
- Outputs are passed on to the world outside the neural network

# Network architecture

## Feedforward networks

- Also known as **multilayer perceptron** (MLP)
- Most widely used architecture
- A connection is allowed from a node in layer  $i$  only to nodes in layer  $i + 1$




# Training algorithms

## Problem statement (I)

ANN can perform different tasks

- Classification, regression, others

Classification (or supervised learning) uses a training set



A <sub>1</sub>	A <sub>2</sub>	A <sub>3</sub>	O	Y
1,1	2,5	4,5	0,2	-0,1
0,9	2,4	1,2	0,5	0,4
1,0	2,0	9,9	0,4	1,2

Toss function: Measure of the error

- Usually mean squared error (mse):  $E = \frac{1}{2}(\gamma - o)^2 = f(w)$
- Y and O are the desired and observed outputs

# Training algorithms

## Problem statement (II)

$$E = \frac{1}{2} \text{Err}^2 = \frac{1}{2} \left[ \gamma - g \left( \sum_{j=0}^n w_j x_j \right) \right]^2$$

where

$\gamma$  Desired output

$w_j$  Weight connection  $j$

$x_j$  Input  $j$

Problem: Determine  $w$  that minimize  $f(w)$

- This is a classical optimization problem
- Any optimization algorithm can be used
- ... in AI, optimization means search

# Training algorithms

## Gradient Descent Algorithm (I)

Given the error

$$E = \frac{1}{2} \text{Err}^2$$

Take partial derivatives

$$\begin{aligned} \frac{\partial E}{\partial w_j} &= \text{Err} \frac{\partial \text{Err}}{\partial w_j} \\ &= \text{Err} \frac{\partial}{\partial w_j} g \left( y - \sum_{j=0}^n w_j x_j \right) \\ &= -\text{Err} \times g'(w) \times x_j \end{aligned}$$

# Training algorithms

## Gradient Descent Algorithm (II)

### Weight update

$$w_j^{k+1} = w_j^k + \alpha \times \text{Err} \times g'(\text{in}) \times x_j$$

with

$\alpha$  Learning rate ( $|\alpha| < 1$ )

err Difference desired and current output

$g'$  Derivate of activation function

$x_j$  Input  $j$

Each iteration is named **epoch**


### Learning algorithm (single neuron)

1. Apply input signal and compute output
2. If output == desired output, do nothing
3. If output < desired output, increase weights
4. If output > desired output, decrease weights




# Training algorithms

## Gradient Descent Algorithm (III)



(Source)



(Source)

# Training algorithms

## Stochastic Gradient Descent (I)

It approximates the gradient by taking samples of the training set

On-line One sample

Mini-batch Several samples

Batch All the samples

Weights update rule:  $w^{k+1} = w^k - \alpha \nabla g(\text{in})$

- where  $\alpha$  is the learning rate

SGD is slow and prone to local minima

# Training algorithms

## Stochastic Gradient Descent (II)

Usually, a momentum is introduced:  $w^{k+1} = w^k - \alpha z^{k+1}$ , where  $z^{k+1} = \beta z^k + \nabla g(\text{in})$

- $\alpha$  is the learning rate
- $\beta$  is the momentum strength
- If  $\beta = 0$  then gradient descend


(On-line demo)

# Training algorithms

## Other optimization algorithms

### Other optimization algorithms

- Newton's method
- Quasi-Newton's method
- Levenberg-Marquardt method
- Conjugate Gradient



# Training algorithms


## Backpropagation algorithm (I)

Efficient learning algorithm for multilayer perceptrons. Three steps

1. **Feed-forward step.** Feed input, compute output and error
2. **Feed-backward step.** Compute individual contribution to error
3. **Adjust weights.** Modify weights to minimize error: Input, output and hidden layers

# Training algorithms

## Backpropagation algorithm (II)



Output layer: Same as single neuron

$$W_j \leftarrow W_j + \alpha \times \text{Err} \times g'(in) \times x_j$$

Define modified error as

$\Delta_i = \text{Err}_i \times g'(in_i)$ , then

$$W_{j,i} \leftarrow W_{j,i} + \alpha \times a_j \times \Delta_i$$

# Training algorithms

## Backpropagation algorithm (III)

Hidden layer: Propagate error

$$W_{k,j} \leftarrow W_{k,j} + \alpha \times a_k \times \Delta_j$$

where

$$\Delta_j = g'(\text{in}) \times \sum_j W_{j,i} \Delta_i$$

## Backpropagation algorithm


1. Compute output
2. Compute output error  $\Delta$
3. For each layer, repeat the following steps
  - 3.1 Propagate Delta backwards
  - 3.2 Update weights between two layers

# Training algorithms

## Learning problems

### Potential problems

- Local minima
- Flat plateau
- Oscillation
- Missing good minima





# Training algorithms

## Learning problems: Under and overfitting (I)

**Underfitting:** Does not learn


- Topology too simple

**Overfitting:** Memorizes samples

- Topology too complex
- Perhaps, the most serious concern in ML
- The net fails when exposed to new data

### Under- and Over-fitting examples


Regression:



predictor too inflexible:  
cannot capture pattern

predictor too flexible:  
fits noise in the data

Classification:



Copyright © 2014 Victor Lomonosov

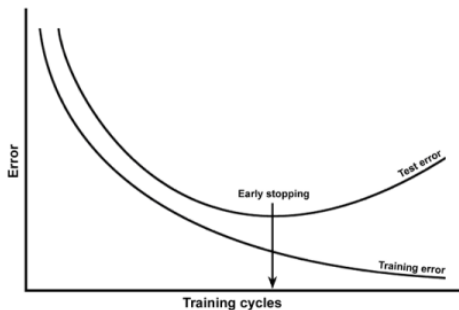
(Source)

# Training algorithms

## Learning problems: Under and overfitting (II)

Solution: Evaluate generalization capabilities

- Split training and validation sets and measure errors



(Source)

# Acknowledgements

- Jesus Aguilar Ruiz, Pablo de Olavide, Seville, Spain
- Daniel Rodríguez, Universidad de Alcalá, Alcalá de Henares, Spain