

Search

Outline

- **Introduction**
- Problem-solving agents
- Problem formulation
- Problem types
- Example problems
- Basic search algorithms
- Conclusions

Introduction

- Early AI works were directed to:
 - Proof of theorems
 - Solving crosswords
 - Games

- All in AI is search
 - Not entirely true (obviously) but more than you can imagine
 - Finding a good/best solution to a problem among several possible solutions


Outline

- ☐ Introduction
- ☐ **Agents**
- ☐ Problem-solving agents
- ☐ Problem formulation
- ☐ Problem types
- ☐ Example problems
- ☐ Basic search algorithms
- ☐ Conclusions

Agents

- An **agent** is anything that can be viewed as **perceiving** its **environment** through **sensors** and **acting** upon that environment through **actuators**
- Human agent: eyes, ears, and other organs for sensors; hands, legs, mouth, and other body parts for actuators
- Robotic agent: cameras and infrared range finders for sensors; various motors for actuators

Agents and environments



- The **agent function** maps from percept histories to actions


$$[f: P^* \rightarrow \mathcal{A}]$$

- The **agent program** runs on the physical **architecture** to produce f
- agent = architecture + program

Agents

- An agent's behavior depends only on its percept sequence (table of action)
 - Mapping from percept sequences to actions
 - Explicit introduce the list
 - Not exhaustively enumerating it by e.g. a function
- If agent's actions are based completely on built-in knowledge (no attention percepts) → lacks autonomy
- Build behaviors based on own experiences and built-in knowledge

Agents: Vacuum-cleaner world



- Percepts: location and contents, e.g., [A, Dirty]
- Actions: *Left*, *Right*, *Suck*, *NoOp*

Vacuum-cleaner agent

Percept sequence	Action
<i>[A, Clean]</i>	<i>Right</i>
<i>[A, Dirty]</i>	<i>Suck</i>
<i>[B, Clean]</i>	<i>Left</i>
<i>[B, Dirty]</i>	<i>Suck</i>
<i>[A, Clean], [A, Clean]</i>	<i>Right</i>
<i>[A, Clean], [A, Dirty]</i>	<i>Suck</i>
<i>⋮</i>	<i>⋮</i>

```
function REFLEX-VACUUM-AGENT( [location,status]) returns an action
  if status = Dirty then return Suck
  else if location = A then return Right
  else if location = B then return Left
```

Outline

- ☐ Introduction
- ☐ Agents
- ☐ **Problem-solving agents**
- ☐ Problem formulation
- ☐ Problem types
- ☐ Example problems
- ☐ Basic search algorithms
- ☐ Conclusions

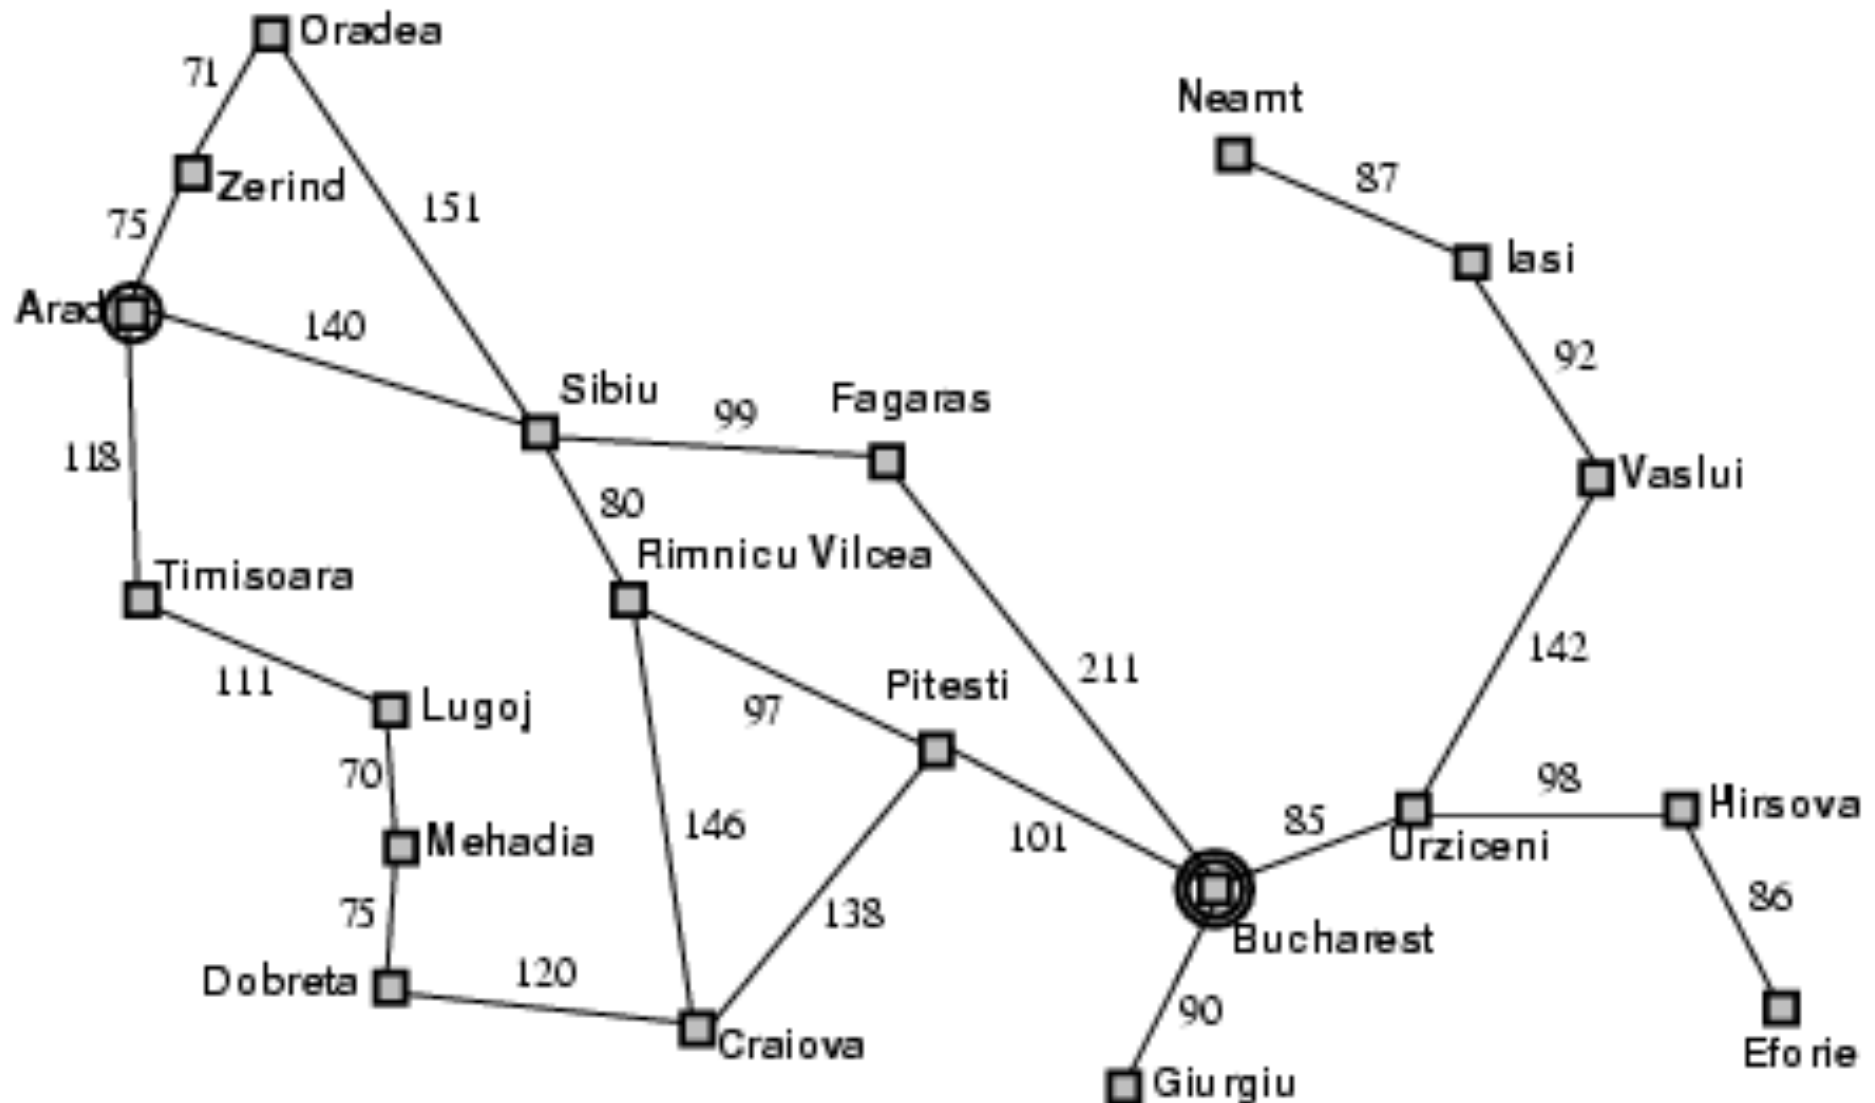
Problem-solving agents (I)

- Agents must maximize its performance measure
- Example: On holiday in Romania; currently in Arad
 - Flight leaves tomorrow from Bucharest
- Formulate goal:
 - be in Bucharest
- Formulating the problem:
 - states: multiple cities
 - actions: drive between cities
- Finding a solution:
 - Sequence cities, eg., Arad, Sibiu, Fagaras, Bucharest
- Process of finding such a solution: **Search**

Problem-solving agents (II)

- Assumptions of the environment:
 - Static: search and formulation is done without considering changes in the environment
 - Observable: the initial state is known
 - Discrete: the alternative locations are known
 - Deterministic: each state is determined by the current state and the action executed
- The solutions are simple sequences of actions, they are executed without considering perceptions

Problem-solving agents (III)



Outline

- ☐ Introduction
- ☐ Agents
- ☐ Problem-solving agents
- ☐ **Problem types**
- ☐ Problem formulation
- ☐ Example problems
- ☐ Basic search algorithms
- ☐ Conclusions

Problem types

- Deterministic, fully observable → **single-state problem**
 - Agent knows exactly which state it will be in; solution is a sequence

- Non-observable → **sensorless problem (conformant problem)**
 - Agent may have no idea where it is; solution is a sequence


- Nondeterministic and/or partially observable → **contingency problem**
 - percepts provide **new** information about current state
 - often **interleave** search with execution

- Unknown state space → **exploration problem**

Example: vacuum world

- **Single-state**, start in #5.


Solution?



Example: vacuum world


- **Single-state**, start in #5.
Solution? [*Right, Suck*]

- **Sensorless**, start in
 $\{1, 2, 3, 4, 5, 6, 7, 8\}$ e.g.,
Right goes to $\{2, 4, 6, 8\}$
Solution?



Example: vacuum world


- **Sensorless**, start in $\{1, 2, 3, 4, 5, 6, 7, 8\}$ e.g.,
Right goes to $\{2, 4, 6, 8\}$
Solution?
[Right, Suck, Left, Suck]



- **Contingency**
 - Nondeterministic: *Suck* may dirty a clean carpet
 - Partially observable: location, dirt at current cell location.
 - Percept: *[L, Clean]*, i.e., start in #5 or #7
Solution?

Example: vacuum world

- **Sensorless**, start in $\{1, 2, 3, 4, 5, 6, 7, 8\}$ e.g.,
Right goes to $\{2, 4, 6, 8\}$
Solution?
[Right, Suck, Left, Suck]



- **Contingency**
 - Nondeterministic: *Suck* may dirty a clean carpet
 - Partially observable: location, dirt at current location.
 - Percept: $[L, \text{Clean}]$, i.e., start in #5 or #7
Solution? *[Right, **if** dirt **then** Suck]*

Outline

- ☐ Introduction
- ☐ Agents
- ☐ Problem-solving agents
- ☐ Problem types
- ☐ **Problem formulation**
- ☐ Example problems
- ☐ Basic search algorithms
- ☐ Conclusions

Problem formulation

- A problem is defined by four items:
 1. **initial state** e.g., "at Arad"
 2. **actions** or **successor function** $S(x)$ = set of action–state pairs
 - e.g., $S(\text{Arad}) = \{ \langle \text{Arad} \rightarrow \text{Zerind}, \text{Zerind} \rangle, \dots \}$
 3. **goal test**, can be
 - **explicit**, e.g., $x = \text{"at Bucharest"}$
 - **implicit**, e.g., $\text{Checkmate}(x)$
 4. **path cost** (additive)
 - e.g., sum of distances, number of actions executed, etc.
 - $c(x, a, y)$ is the step cost, assumed to be ≥ 0
- A **solution** is a sequence of actions leading from the initial state to a goal state


Problem formulation

- Real world is absurdly complex
 - state space must be **abstracted** for problem solving
- (Abstract) state = set of real states
- (Abstract) action = complex combination of real actions
 - e.g., "Arad → Zerind" represents a complex set of possible routes, detours, rest stops, etc.
- For guaranteed realizability, **any** real state "in Arad" must get to **some** real state "in Zerind"
- (Abstract) solution =
 - set of real paths that are solutions in the real world
- Each abstract action should be "easier" than the original problem

Outline


- ☐ Introduction
- ☐ Agents
- ☐ Problem-solving agents
- ☐ Problem types
- ☐ Problem formulation
- ☐ **Example problems**
- ☐ Basic search algorithms
- ☐ Conclusions

Vacuum world state space graph



- ☐ states?
- ☐ actions?
- ☐ goal test?
- ☐ path cost?

Vacuum world state space graph



- states? integer dirt and robot location
- actions? *Left, Right, Suck*
- goal test? no dirt at all locations
- path cost? 1 per action

Example: The 8-puzzle

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State

- ☐ states?
- ☐ actions?
- ☐ goal test?
- ☐ path cost?

Example: The 8-puzzle

7	2	4
5		6
8	3	1

Start State


	1	2
3	4	5
6	7	8

Goal State

- states? locations of tiles
- actions? move blank left, right, up, down
- goal test? = goal state (given)
- path cost? 1 per move


[Note: optimal solution of n -Puzzle family is NP-hard]

Example: robotic assembly



- ☐ states?
- ☐ actions?
- ☐ goal test?
- ☐ path cost?

Example: robotic assembly



- states?: real-valued coordinates of robot joint angles
parts of the object to be assembled
- actions?: continuous motions of robot joints
- goal test?: complete assembly
- path cost?: time to execute



Outline

- ☐ Introduction
- ☐ Agents
- ☐ Problem-solving agents
- ☐ Problem types
- ☐ Problem formulation
- ☐ Example problems
- ☐ **Basic search algorithms**
- ☐ Conclusions


Search algorithms

- We have formulated problems, we now need to solve them: search tree
- In general we can have a search graph rather than a tree when the state can be reached from multiple paths
- Basic idea:
 - offline, simulated exploration of state space by generating successors of already-explored states (a.k.a.~**expanding** states)


Tree search example



Tree search example




Tree search example



Implementation: states vs. nodes

- A **state** is a (representation of) a physical configuration
- A **node** is a data structure constituting part of a search tree includes **state**, **parent node**, **action**, **path cost** $g(x)$, **depth**



- The `Expand` function creates new nodes, filling in the various fields and using the `SuccessorFn` of the problem to create the corresponding states.

Search strategies

- A search strategy is defined by picking the **order of node expansion**
- Strategies are evaluated along the following dimensions:
 - **completeness**: does it always find a solution if one exists?
 - **time complexity**: number of nodes generated
 - **space complexity**: maximum number of nodes in memory
 - **optimality**: does it always find a least-cost solution?
- Time and space complexity are measured in terms of
 - b : maximum branching factor of the search tree
 - d : depth of the least-cost solution
 - m : maximum depth of the state space (may be ∞)

Outline

- Introduction
- Agents
- Problem-solving agents
- Problem types
- Problem formulation
- Example problems
- Basic search algorithms
 - Uninformed search strategies
 - Informed search strategies
- Conclusions

Outline

- Introduction
- Agents
- Problem-solving agents
- Problem types
- Problem formulation
- Example problems
- Basic search algorithms
 - Uninformed search strategies
 - Informed search strategies
- Conclusions

Outline

- ☐ Introduction
- ☐ Agents
- ☐ Problem-solving agents
- ☐ Problem formulation
- ☐ Problem types
- ☐ Example problems
- ☐ Basic search algorithms
- ☐ **Conclusions**

Conclusions

- Problem formulation usually requires abstracting away real-world details to define a state space that can feasibly be explored
- Variety of uninformed and informed search strategies