

# Search

---

# Outline

---

- ☐ Introduction
- ☐ Problem-solving agents
- ☐ Problem formulation
- ☐ Problem types
- ☐ Example problems
- ☐ Basic search algorithms
- ☐ Conclusions

# Introduction

---

- Early AI works were directed to:
  - Proof of theorems
  - Solving crosswords
  - Games
- All in AI is search
  - Not entirely true (obviously) but more than you can imagine
  - Finding a good/best solution to a problem among several possible solutions

# Outline

---

- ☐ Introduction
- ☐ **Problem-solving agents**
- ☐ Problem formulation
- ☐ Problem types
- ☐ Example problems
- ☐ Basic search algorithms
- ☐ Conclusions

# Problem-solving agents

---

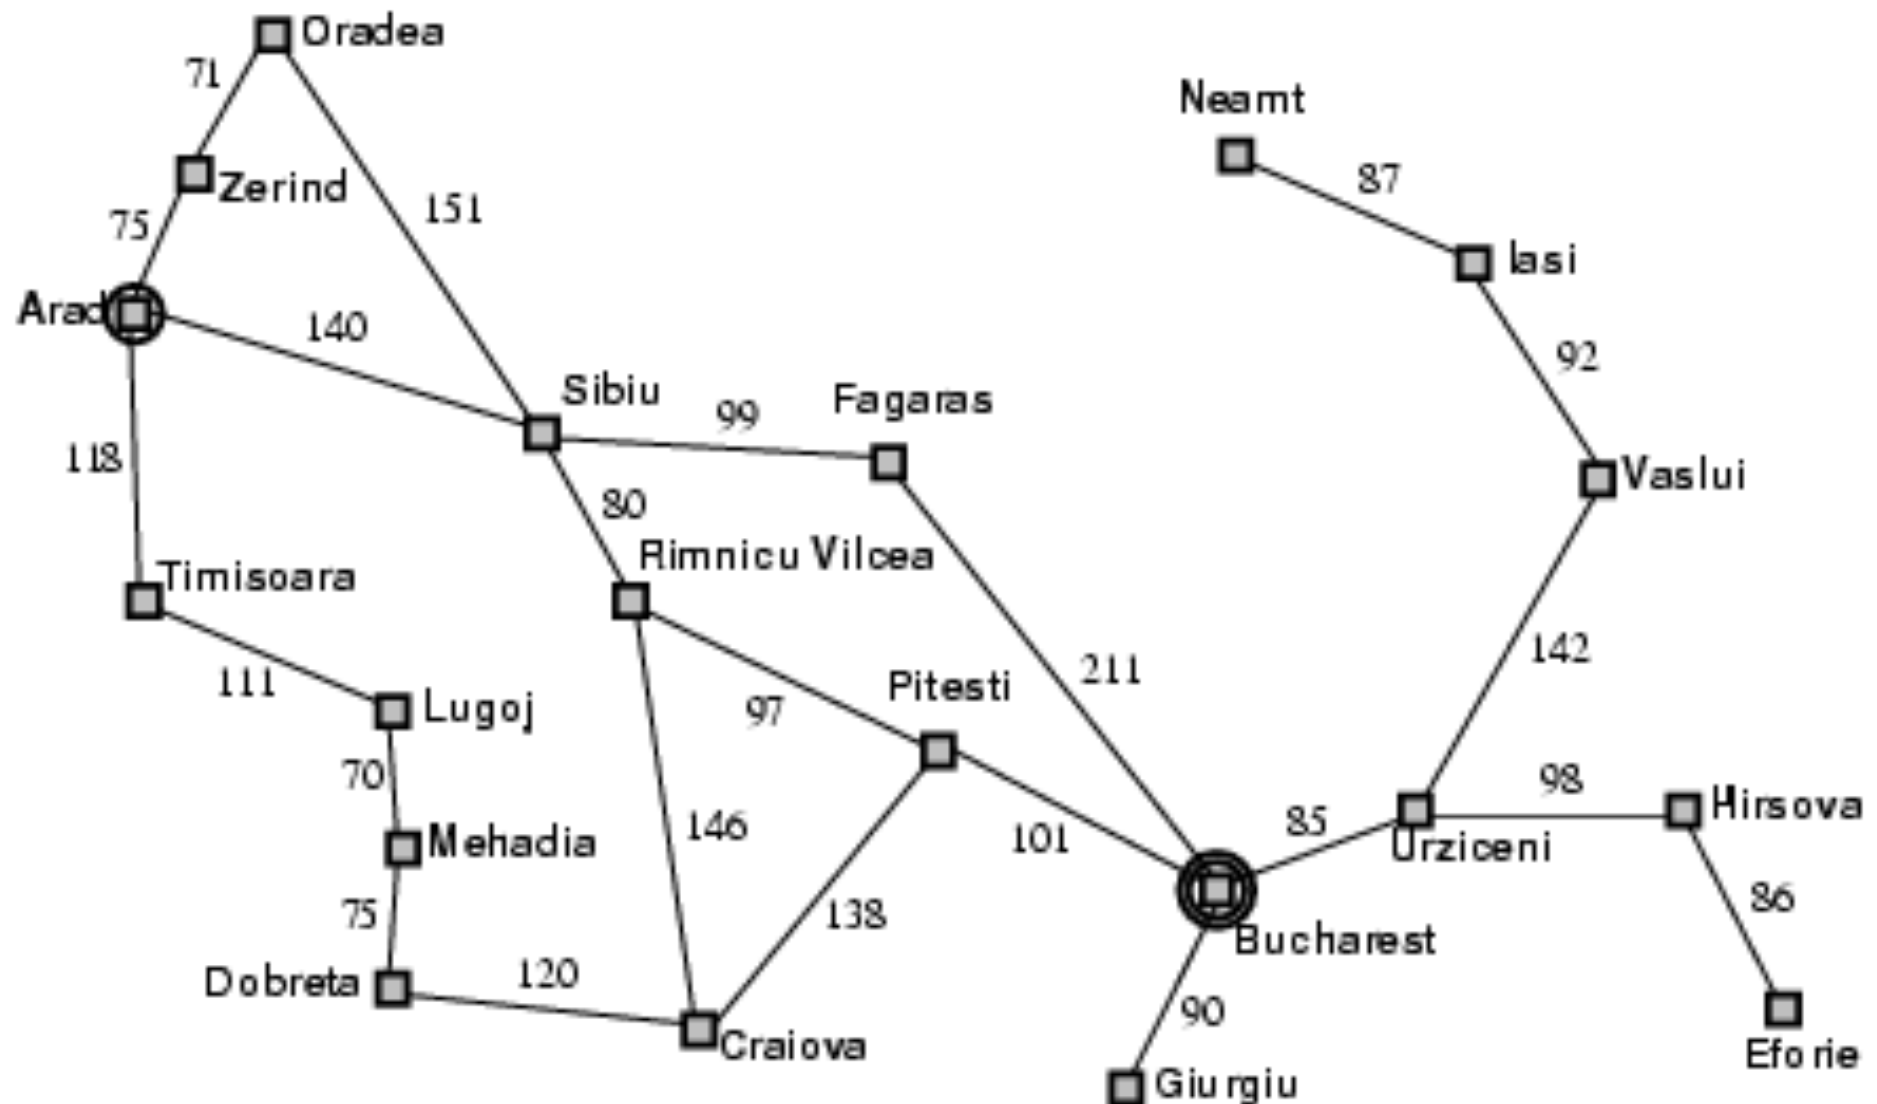
- Agents must maximize its performance measure
- Example: On holiday in Romania; currently in Arad
  - Flight leaves tomorrow from Bucharest
- Formulate goal:
  - be in Bucharest
- Formulating the problem:
  - states: multiple cities
  - actions: drive between cities
- Finding a solution:
  - Sequence cities, eg., Arad, Sibiu, Fagaras, Bucharest
- The process of finding such a solution is called search

# Problem-solving agents

---

- Assumptions of the environment:
  - Static: search and formulation is done without considering changes in the environment
  - Observable: the initial state is known
  - Discrete: the alternative locations are known
  - Deterministic: each state is determined by the current state and the action executed
- The solutions are simple sequences of actions, they are executed without considering perceptions

# Problem-solving agents



# Outline

---

- ☐ Introduction
- ☐ Problem-solving agents
- ☐ **Problem types**
- ☐ Problem formulation
- ☐ Example problems
- ☐ Basic search algorithms
- ☐ Conclusions



# Problem types

---


- Deterministic, fully observable → **single-state problem**
  - Agent knows exactly which state it will be in; solution is a sequence
  
- Non-observable → **sensorless problem (conformant problem)**
  - Agent may have no idea where it is; solution is a sequence
  
- Nondeterministic and/or partially observable → **contingency problem**
  - percepts provide **new** information about current state
  - often **interleave** search with execution
  
- Unknown state space → **exploration problem**

# Example: vacuum world

---

- **Single-state**, start in #5.

Solution?




# Example: vacuum world

---


- **Single-state**, start in #5.  
Solution? [*Right, Suck*]

- **Sensorless**, start in  
 $\{1, 2, 3, 4, 5, 6, 7, 8\}$  e.g.,  
*Right* goes to  $\{2, 4, 6, 8\}$   
Solution?



# Example: vacuum world


- **Sensorless**, start in  $\{1, 2, 3, 4, 5, 6, 7, 8\}$  e.g.,  
*Right* goes to  $\{2, 4, 6, 8\}$   
Solution?  
*[Right, Suck, Left, Suck]*



- **Contingency**
  - Nondeterministic: *Suck* may dirty a clean carpet
  - Partially observable: location, dirt at current location.
  - Percept:  $[L, \text{Clean}]$ , i.e., start in #5 or #7Solution?

# Example: vacuum world

- **Sensorless**, start in  $\{1, 2, 3, 4, 5, 6, 7, 8\}$  e.g.,  
*Right* goes to  $\{2, 4, 6, 8\}$   
Solution?  
*[Right, Suck, Left, Suck]*



- **Contingency**
  - Nondeterministic: *Suck* may dirty a clean carpet
  - Partially observable: location, dirt at current location.
  - Percept: *[L, Clean]*, i.e., start in #5 or #7  
Solution? *[Right, **if** dirt **then** Suck]*

# Outline

---

- ☐ Introduction
- ☐ Problem-solving agents
- ☐ Problem types
- ☐ **Problem formulation**
- ☐ Example problems
- ☐ Basic search algorithms
- ☐ Conclusions

# Problem formulation

---

- A problem is defined by four items:
  1. **initial state** e.g., "at Arad"
  2. **actions** or **successor function**  $S(x)$  = set of action–state pairs
    - e.g.,  $S(\text{Arad}) = \{ \langle \text{Arad} \rightarrow \text{Zerind}, \text{Zerind} \rangle, \dots \}$
  3. **goal test**, can be
    - **explicit**, e.g.,  $x = \text{"at Bucharest"}$
    - **implicit**, e.g.,  $\text{Checkmate}(x)$
  4. **path cost** (additive)
    - e.g., sum of distances, number of actions executed, etc.
    - $c(x, a, y)$  is the step cost, assumed to be  $\geq 0$
- A **solution** is a sequence of actions leading from the initial state to a goal state

# Problem formulation

---

- Real world is absurdly complex
  - state space must be **abstracted** for problem solving
- (Abstract) state = set of real states
- (Abstract) action = complex combination of real actions
  - e.g., "Arad → Zerind" represents a complex set of possible routes, detours, rest stops, etc.
- For guaranteed realizability, **any** real state "in Arad" must get to **some** real state "in Zerind"
- (Abstract) solution =
  - set of real paths that are solutions in the real world
- Each abstract action should be "easier" than the original problem




# Outline

---


- ☐ Introduction
- ☐ Problem-solving agents
- ☐ Problem types
- ☐ Problem formulation
- ☐ **Example problems**
- ☐ Basic search algorithms
- ☐ Conclusions

# Vacuum world state space graph



- ☐ states?
- ☐ actions?
- ☐ goal test?
- ☐ path cost?

# Vacuum world state space graph



- states? integer dirt and robot location
- actions? *Left, Right, Suck*
- goal test? no dirt at all locations
- path cost? 1 per action

# Example: The 8-puzzle

---

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State

- ☐ states?
- ☐ actions?
- ☐ goal test?
- ☐ path cost?

# Example: The 8-puzzle

---

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8


Goal State

- states? locations of tiles
- actions? move blank left, right, up, down
- goal test? = goal state (given)
- path cost? 1 per move

[Note: optimal solution of  $n$ -Puzzle family is NP-hard]

# Example: robotic assembly


---



- ☐ states?
- ☐ actions?
- ☐ goal test?
- ☐ path cost?

# Example: robotic assembly

---



- states?: real-valued coordinates of robot joint angles  
parts of the object to be assembled
- actions?: continuous motions of robot joints
- goal test?: complete assembly
- path cost?: time to execute

# Outline

---

- ☐ Introduction
- ☐ Problem-solving agents
- ☐ Problem types
- ☐ Problem formulation
- ☐ Example problems
- ☐ **Basic search algorithms**
- ☐ Conclusions

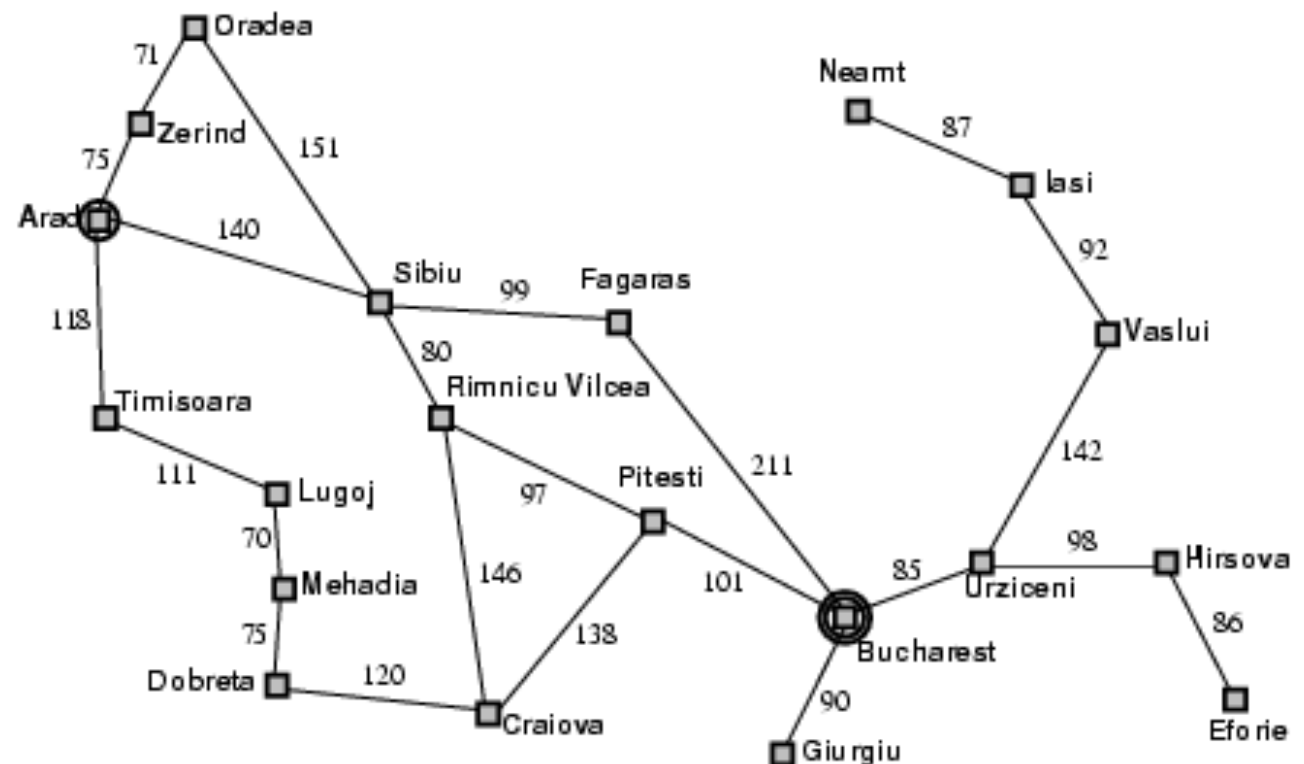



# Search algorithms

---


- We have formulated problems, we now need to solve them: search tree
- In general we can have a search graph rather than a tree when the state can be reached from multiple paths
- Basic idea:
  - offline, simulated exploration of state space by generating successors of already-explored states (a.k.a. ~**expanding** states)

# Tree search example




# Tree search example

---




# Tree search example

---



# Implementation: states vs. nodes

- A **state** is a (representation of) a physical configuration
- A **node** is a data structure constituting part of a search tree includes **state**, **parent node**, **action**, **path cost**  $g(x)$ , **depth**



- The `Expand` function creates new nodes, filling in the various fields and using the `SuccessorFn` of the problem to create the corresponding states.

# Search strategies

---

- A search strategy is defined by picking the **order of node expansion**
- Strategies are evaluated along the following dimensions:
  - **completeness**: does it always find a solution if one exists?
  - **time complexity**: number of nodes generated
  - **space complexity**: maximum number of nodes in memory
  - **optimality**: does it always find a least-cost solution?
- Time and space complexity are measured in terms of
  - $b$ : maximum branching factor of the search tree
  - $d$ : depth of the least-cost solution
  - $m$ : maximum depth of the state space (may be  $\infty$ )

# Outline

---

- Introduction
- Problem-solving agents
- Problem types
- Problem formulation
- Example problems
- Basic search algorithms
  - **Uninformed search strategies**
  - Informed search strategies
- Conclusions

# Uninformed search strategies

---


- **Uninformed** search strategies use only the information available in the problem definition
  - Breadth-first search
  - Uniform-cost search
  - Depth-first search
  - Depth-limited search
  - Iterative deepening search



# Breadth-first search

---


- Expand shallowest unexpanded node
- Implementation:
  - *fringe* is a FIFO queue, i.e., new successors go at end



# Breadth-first search

---


- Expand shallowest unexpanded node
- Implementation:
  - *fringe* is a FIFO queue, i.e., new successors go at end



# Breadth-first search

---


- Expand shallowest unexpanded node
- Implementation:
  - *fringe* is a FIFO queue, i.e., new successors go at end



# Breadth-first search

---

- Expand shallowest unexpanded node
- Implementation:
  - *fringe* is a FIFO queue, i.e., new successors go at end



# Properties of breadth-first search

---

- ❑ Complete? Yes (if  $b$  is finite)
- ❑ Time?  $1+b+b^2+b^3+\dots +b^d + b(b^d-1) = O(b^{d+1})$
- ❑ Space?  $O(b^{d+1})$  (keeps every node in memory)
- ❑ Optimal? Yes (if cost = 1 per step)

**Space** is the bigger problem (more than time)

Each state has  $b$  successors (branching factor)  
 $d$  is the shallower depth

# Uniform-cost search

---


- ❑ Expand least-cost unexpanded node
- ❑ **Implementation:**
  - *fringe* = queue ordered by path cost
- ❑ Equivalent to breadth-first if step costs all equal
- ❑ Complete? Yes, if step cost  $\geq \epsilon$
- ❑ Time? # of nodes with  $g \leq$  cost of optimal solution,  $O(b^{\text{ceiling}(C^*/\epsilon)})$  where  $C^*$  is the cost of the optimal solution
- ❑ Space? # of nodes with  $g \leq$  cost of optimal solution,  $O(b^{\text{ceiling}(C^*/\epsilon)})$
- ❑ Optimal? Yes – nodes expanded in increasing order of  $g(n)$

*If all costs are equal  $\rightarrow O(b^d)$*

# Depth-first search

---


- Expand deepest unexpanded node
- **Implementation:**
  - *fringe* = LIFO queue, i.e., put successors at front



# Depth-first search

---

- Expand deepest unexpanded node
- **Implementation:**
  - *fringe* = LIFO queue, i.e., put successors at front






# Depth-first search

---


- Expand deepest unexpanded node
- **Implementation:**
  - *fringe* = LIFO queue, i.e., put successors at front



# Depth-first search

---


- Expand deepest unexpanded node
- **Implementation:**
  - *fringe* = LIFO queue, i.e., put successors at front



# Depth-first search

---


- Expand deepest unexpanded node
- **Implementation:**
  - *fringe* = LIFO queue, i.e., put successors at front



# Depth-first search

---


- Expand deepest unexpanded node
- **Implementation:**
  - *fringe* = LIFO queue, i.e., put successors at front



# Depth-first search

---


- Expand deepest unexpanded node
- **Implementation:**
  - *fringe* = LIFO queue, i.e., put successors at front



# Depth-first search

---


- Expand deepest unexpanded node
- **Implementation:**
  - *fringe* = LIFO queue, i.e., put successors at front



# Depth-first search

---


- Expand deepest unexpanded node
- **Implementation:**
  - *fringe* = LIFO queue, i.e., put successors at front



# Depth-first search

---

- Expand deepest unexpanded node
- **Implementation:**
  - *fringe* = LIFO queue, i.e., put successors at front






# Depth-first search

---


- Expand deepest unexpanded node
- **Implementation:**
  - *fringe* = LIFO queue, i.e., put successors at front



# Depth-first search

---

- Expand deepest unexpanded node
- **Implementation:**
  - *fringe* = LIFO queue, i.e., put successors at front



# Properties of depth-first search

---

- Complete? No: fails in infinite-depth spaces, spaces with loops
  - Modify to avoid repeated states along path  
→ complete in finite spaces
- Time?  $O(b^m)$ : terrible if  $m$  is much larger than  $d$ 
  - but if solutions are dense, may be much faster than breadth-first
- Space?  $O(bm)$
- Optimal? No

# Depth-limited search

---

□ = depth-first search with depth limit  $l$ ,  
i.e., nodes at depth  $l$  have no successor

# Iterative deepening search / =0

---



Limit = 0



# Iterative deepening search / =1


---

Limit = 1




# Iterative deepening search / =2

Limit = 2



# Iterative deepening search / =3

Limit = 3





# Iterative deepening search

---

- Number of nodes generated in a depth-limited search to depth  $d$  with branching factor  $b$ :

$$N_{DLS} = b^0 + b^1 + b^2 + \dots + b^{d-2} + b^{d-1} + b^d$$

- Number of nodes generated in an iterative deepening search to depth  $d$  with branching factor  $b$ :

$$N_{IDS} = (d+1)b^0 + d b^1 + (d-1)b^2 + \dots + 3b^{d-2} + 2b^{d-1} + 1b^d$$

- For  $b = 10$ ,  $d = 5$ ,

- $N_{DLS} = 1 + 10 + 100 + 1,000 + 10,000 + 100,000 = 111,111$

- $N_{IDS} = 6 + 50 + 400 + 3,000 + 20,000 + 100,000 = 123,456$

- Overhead =  $(123,456 - 111,111)/111,111 = 11\%$

# Properties of iterative deepening search

---

- Complete? Yes
- Time?  $(d+1)b^0 + d b^1 + (d-1)b^2 + \dots + b^d = O(b^d)$
- Space?  $O(bd)$
- Optimal? Yes, if step cost = 1

# Summary of algorithms



---

Criterion	Breadth-First	Uniform-Cost	Depth-First	Depth-Limited	Iterative Deepening
Complete?	Yes	Yes	No	No	Yes
Time	$O(b^{d+1})$	$O(b^{\lceil C^*/\epsilon \rceil})$	$O(b^m)$	$O(b^l)$	$O(b^d)$
Space	$O(b^{d+1})$	$O(b^{\lceil C^*/\epsilon \rceil})$	$O(bm)$	$O(bl)$	$O(bd)$
Optimal?	Yes	Yes	No	No	Yes

# Repeated states

---

- ❑ Failure to detect repeated states can turn a linear problem into an exponential one!



# Outline

---

- Introduction
- Problem-solving agents
- Problem types
- Problem formulation
- Example problems
- Basic search algorithms
  - Uninformed search strategies
  - **Informed search strategies**
- Conclusions

# Informed search strategies

---

- Use the problem-specific knowledge beyond the definition of the problem itself to find more efficient solutions than uninformed strategy
  - Best-first search
    - Greedy best-first search
    - A\* search
  - Heuristics
  - Local search algorithms
    - Hill-climbing search
    - Simulated annealing search
    - Local beam search
    - Genetic algorithms


# Best-first search

---

- Idea: use an **evaluation function**  $f(n)$  for each node
  - estimate of "desirability"
  - Expand most desirable unexpanded node
  
- Implementation:

Order the nodes in fringe in decreasing order of desirability
  
- Special cases:
  - greedy best-first search
  - A\* search

# Romania with step costs in km





# Informed search strategies

---

- Best-first search
  - **Greedy best-first search**
  - $A^*$  search
- Heuristics
- Local search algorithms
  - Hill-climbing search
  - Simulated annealing search
  - Local beam search
  - Genetic algorithms

# Greedy best-first search

---


- Evaluation function  $f(n) = h(n)$  (**h**euristic) = estimate of cost from  $n$  to *goal*
- Greedy best-first search expands the node that **appears** to be closest to the goal
- Implementation: as a priority queue to keep the fringe in ascending order of  $f$ -values
- e.g.,  $h_{SLD}(n)$  = straight-line distance from  $n$  to Bucharest

# Greedy best-first search example

---




# Greedy best-first search example



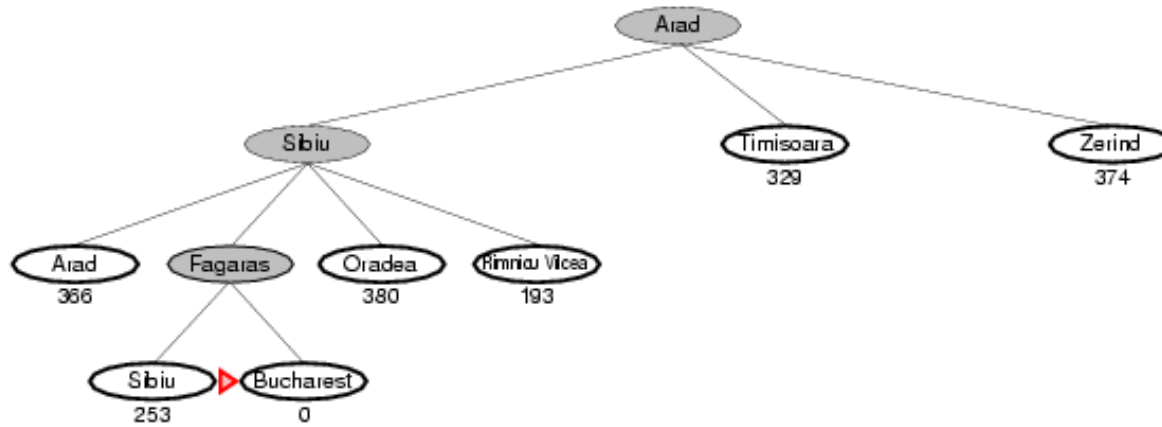
# Greedy best-first search example

---



# Greedy best-first search example

---



# Properties of greedy best-first search

---

- ❑ Complete? No – can get stuck in loops, e.g., lasi  
→ Neamt → lasi → Neamt →
- ❑ Time?  $O(b^m)$ , but a good heuristic can give dramatic improvement
- ❑ Space?  $O(b^m)$  -- keeps all nodes in memory
- ❑ Optimal? No

Similar to depth-first search

Each state has  $b$  successors (branching factor)

$d$  is the depth of the shallowest solution

# Informed search strategies

---

- Best-first search
  - Greedy best-first search
  - **A\* search**
- Heuristics
- Local search algorithms
  - Hill-climbing search
  - Simulated annealing search
  - Local beam search
  - Genetic algorithms



# A\* search

---

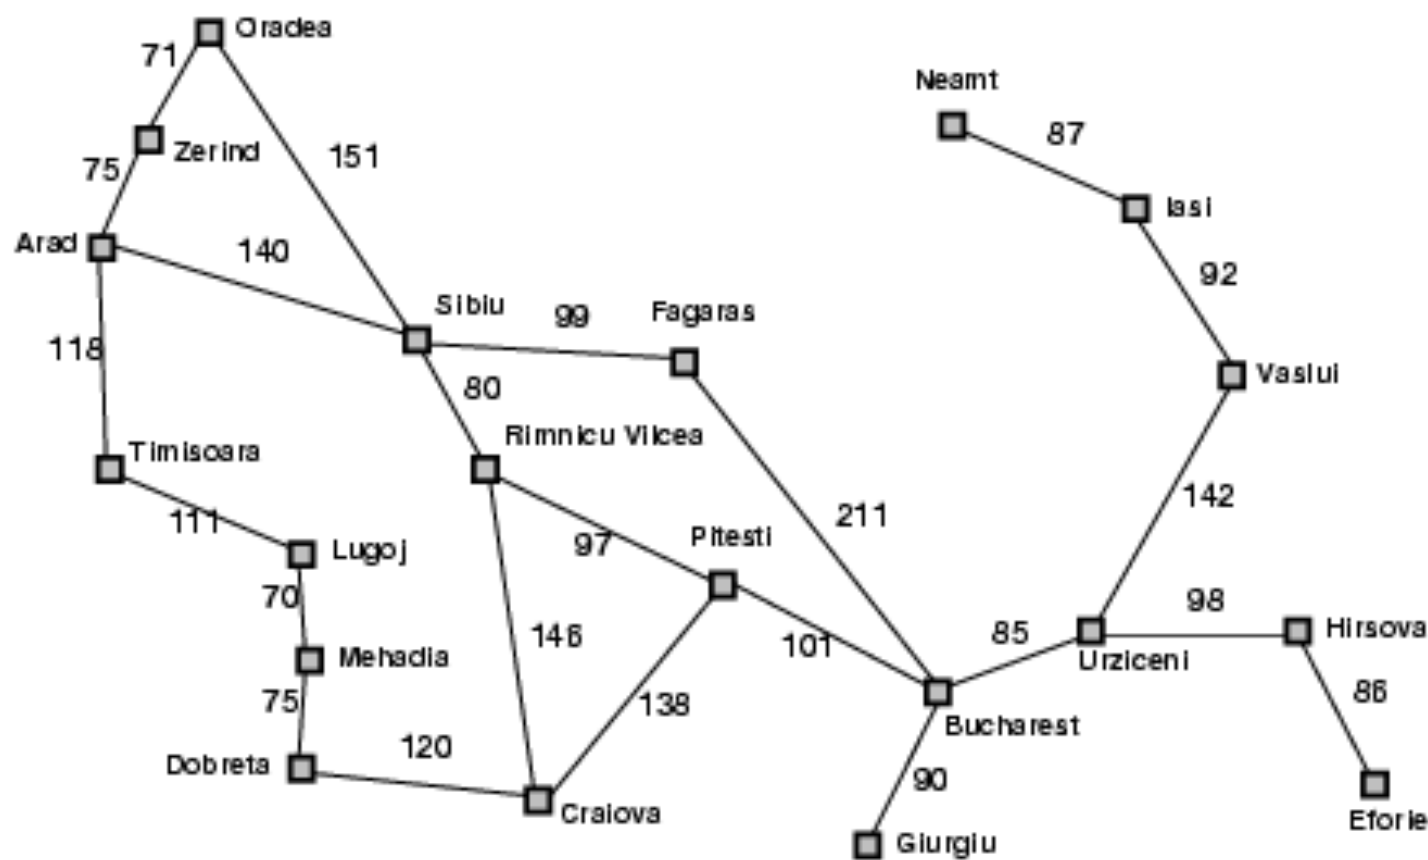
- Idea: avoid expanding paths that are already expensive
- Evaluation function  $f(n) = g(n) + h(n)$ 
  - $g(n)$  = cost so far to reach  $n$
  - $h(n)$  = estimated cost from  $n$  to goal
  - $f(n)$  = estimated total cost of path through  $n$  to goal
- A\* is optimal if  $h(n)$  is an admissible heuristic such that  $h(n)$  never overestimates the cost to reach the goal

# A\* search example

---

▶ Arad  
366=0+366



# A\* search example



Straight-line distance  
to Bucharest

<b>Arad</b>	366
<b>Bucharest</b>	0
<b>Craiova</b>	160
<b>Dobreta</b>	242
<b>Eforie</b>	161
<b>Fagaras</b>	176
<b>Giurgiu</b>	77
<b>Hirsova</b>	151
<b>Iasi</b>	226
<b>Lugoj</b>	244
<b>Mehadia</b>	241
<b>Neamt</b>	234
<b>Oradea</b>	380
<b>Pitesti</b>	100
<b>Rimnicu Vilcea</b>	193
<b>Sibiu</b>	253
<b>Timisoara</b>	329
<b>Urziceni</b>	80
<b>Vaslui</b>	199
<b>Zerind</b>	374

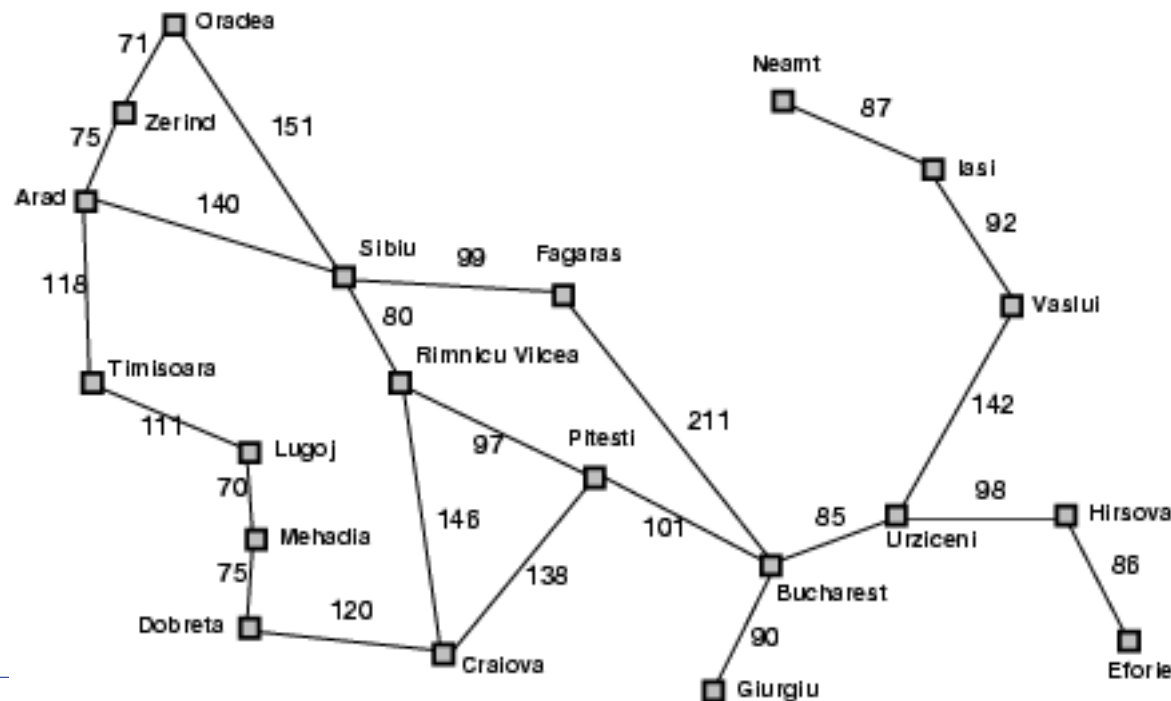
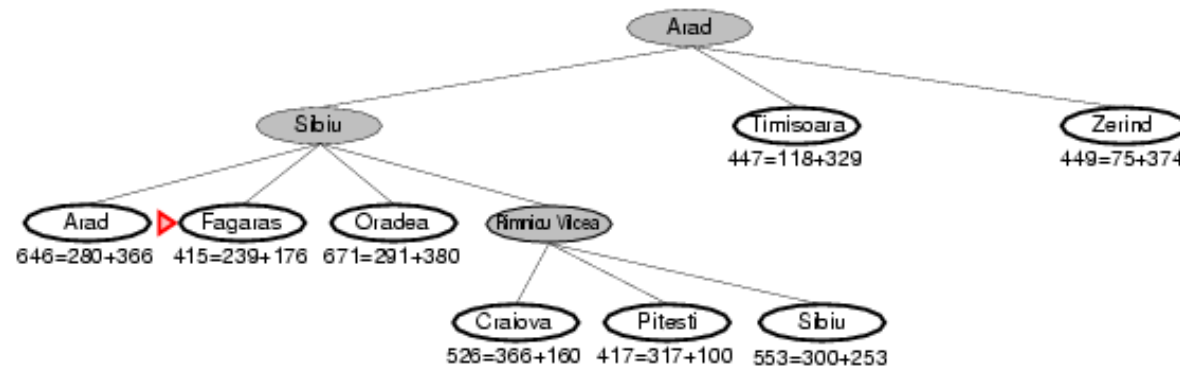
# A\* search example



Straight-line distance  
to Bucharest

Arad	366
Bucharest	0
Craiova	160
Dobreta	242
Eforie	161
Fagaras	176
Giurgiu	77
Hirsova	151
Iasi	226
Lugoj	244
Mehadia	241
Neamt	234
Oradea	380
Pitesti	10
Rimnicu Vilcea	193
Sibiu	253
Timisoara	329
Urziceni	80
Vaslui	199
Zerind	374

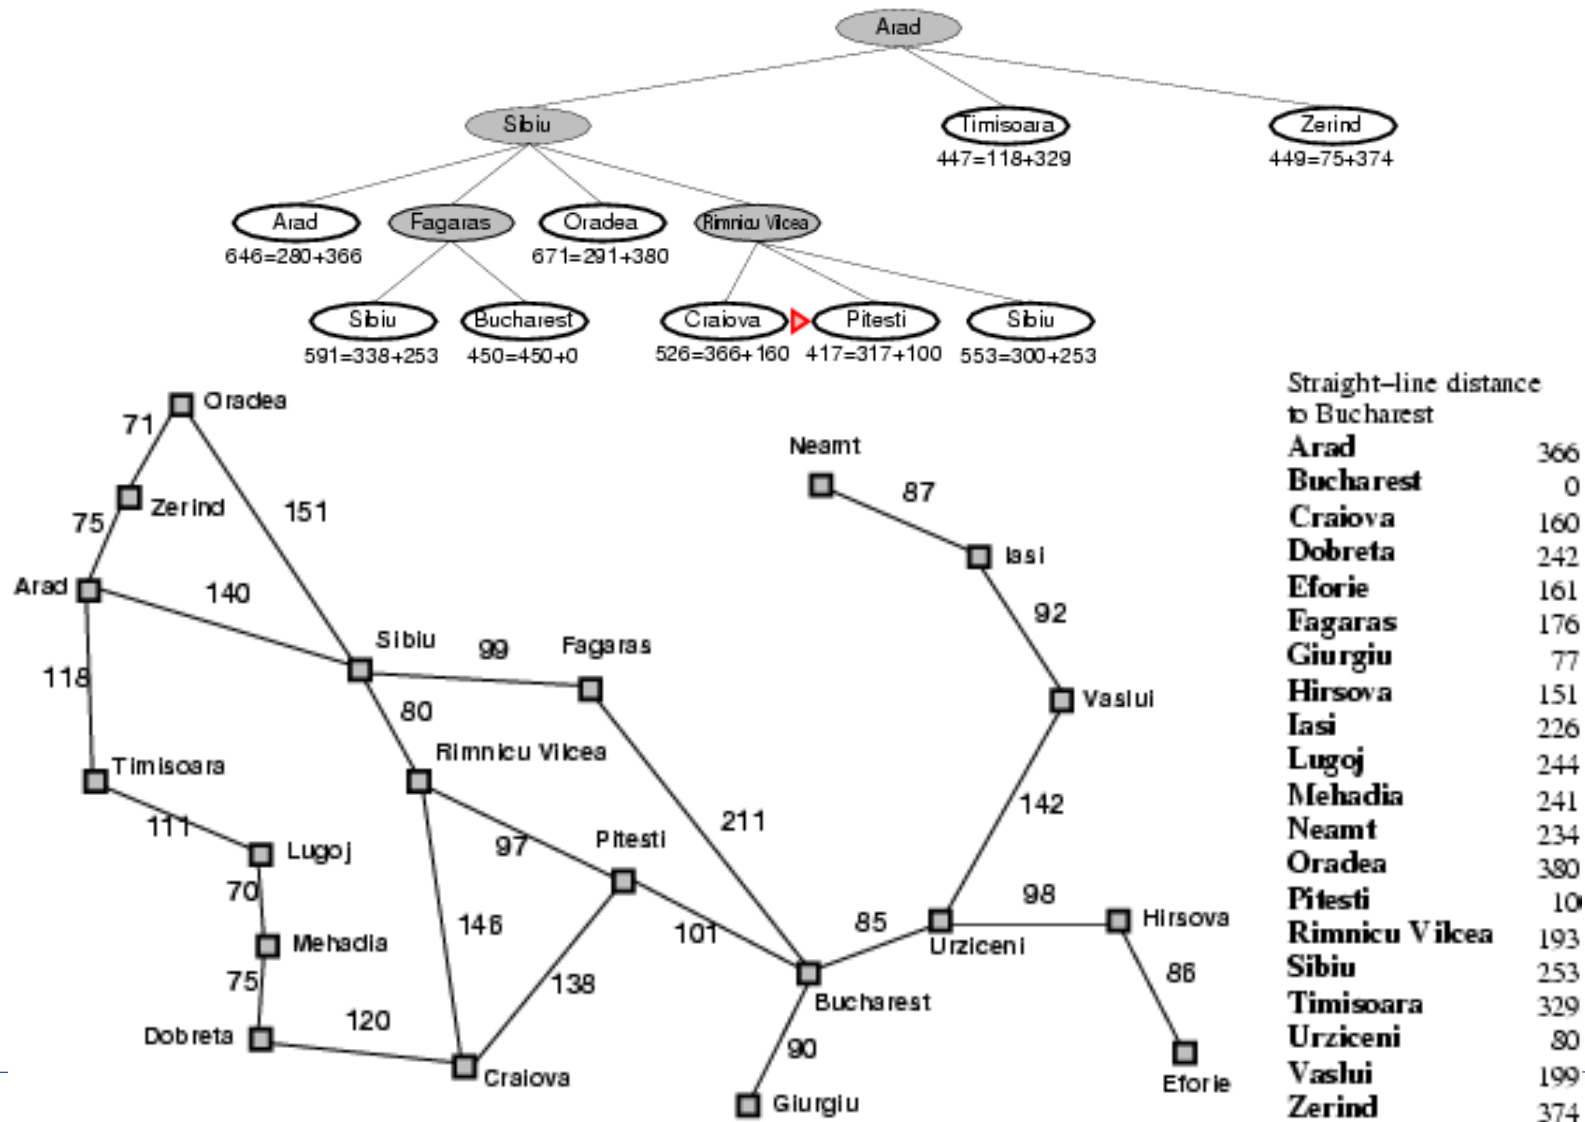
# A\* search example




Straight-line distance  
to Bucharest

Arad	366
Bucharest	0
Craiova	160
Dobreta	242
Eforie	161
Fagaras	176
Giurgiu	77
Hirsova	151
Iasi	226
Lugoj	244
Mehadia	241
Neamt	234
Oradea	380
Pitesti	10
Rimnicu Vilcea	193
Sibiu	253
Timisoara	329
Urziceni	80
Vaslui	199
Zerind	374

# A\* search example



# A\* search example



# Informed search strategies

---

- ☐ Best-first search
- ☐ Greedy best-first search
- ☐ A\* search
- ☐ **Heuristics**
- ☐ Local search algorithms
  - Hill-climbing search
  - Simulated annealing search
  - Local beam search
  - Genetic algorithms



# Admissible heuristics

---

- A heuristic  $h(n)$  is **admissible** if for every node  $n$ ,  $h(n) \leq h^*(n)$ , where  $h^*(n)$  is the **true** cost to reach the goal state from  $n$
- An admissible heuristic **never overestimates** the cost to reach the goal, i.e., it is **optimistic**
- Example:  $h_{SLD}(n)$  (never overestimates the actual road distance)
- **Theorem**: If  $h(n)$  is admissible,  $A^*$  using TREE-SEARCH is optimal

# Properties of A\*

---

- ❑ Complete? Yes (unless there are infinitely many nodes with  $f \leq f(G)$  )
- ❑ Time? Exponential
- ❑ Space? Keeps all nodes in memory
- ❑ Optimal? Yes

# Admissible heuristics for 8-puzzle

---

- ❑ The average cost for the 8-puzzle are approx. 22 steps. Here are 26 steps.
- ❑ Branching factor is approx. 3
  - Empty in the middle, 4 mov
  - Empty in the corner, 2 mov
  - Rest cases, 3 mov
- ❑ Depth-first search will look  $3^{22}$  states
- ❑ If we keep track of repeated states, we could reduce to 170.000
- ❑ In the 15-puzzle =  $10^{13}$

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State

# Admissible heuristics for 8-puzzle

E.g., for the 8-puzzle:

- $h_1(n)$  = number of misplaced tiles
- $h_2(n)$  = total Manhattan distance  
(i.e., + horizontal and vertical distance from desired location of each tile)

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State

- $h_1(S) = ?$
- $h_2(S) = ?$

# Admissible heuristics

E.g., for the 8-puzzle:

- $h_1(n)$  = number of misplaced tiles
- $h_2(n)$  = total Manhattan distance  
(i.e., no. of squares from desired location of each tile)

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State

- $h_1(S) = ?$  8
- $h_2(S) = ?$   $3+1+2+2+2+3+3+2 = 18$

# Dominance

---

- If  $h_2(n) \geq h_1(n)$  for all  $n$  (both admissible) then  $h_2$  **dominates**  $h_1$
- $h_2$  is better for search
- Typical search costs (average number of nodes expanded):
  - $d=12$       IDS = 3,644,035 nodes  
           $A^*(h_1) = 227$  nodes  
           $A^*(h_2) = 73$  nodes
  - $d=24$       IDS = too many nodes  
           $A^*(h_1) = 39,135$  nodes  
           $A^*(h_2) = 1,641$  nodes

# Relaxed problems

---

- A problem with fewer restrictions on the actions is called a **relaxed problem**
- The cost of an optimal solution to a relaxed problem is an admissible heuristic for the original problem
- If the rules of the 8-puzzle are relaxed so that a tile can move **anywhere**, then  $h_1(n)$  gives the shortest solution
- If the rules are relaxed so that a tile can move to **any adjacent square**, then  $h_2(n)$  gives the shortest solution

# Informed search strategies

---

- ☐ Best-first search
- ☐ Greedy best-first search
- ☐ A\* search
- ☐ Heuristics
- ☐ **Local search algorithms**
  - Hill-climbing search
  - Simulated annealing search
  - Local beam search
  - Genetic algorithms



# Local search algorithms

---

- In many optimization problems, the **path** to the goal is irrelevant; the goal state itself is the solution
- State space = set of "complete" configurations. Find configuration satisfying constraints, e.g., n-queens
- In such cases, we can use **local search algorithms**
- Keep a single "current" state, try to improve it
- Work with one current state and generally moves to the neighboring state
- The paths followed by the search are not retained
  - They use little memory
  - You can find reasonable solutions in large state spaces or infinite

# Example: $n$ -queens

---

- Put  $n$  queens on an  $n \times n$  board with no two queens on the same row, column, or diagonal



# Informed search strategies

---

- ☐ Best-first search
- ☐ Greedy best-first search
- ☐ A\* search
- ☐ Heuristics
- ☐ Local search algorithms
  - **Hill-climbing search**
  - Simulated annealing search
  - Local beam search
  - Genetic algorithms

# Hill-climbing search


---

- It's just a loop that moves in the direction of increasing value
  - Ends when it reaches a peak where no neighbor has a higher value
  - The search tree is not kept, just a data structure of the current node to check the goal condition and its objective function value
- "Like climbing Everest in thick fog with amnesia"

# Hill-climbing search


---

- Problem: depending on initial state, can get stuck in local maxima



# Hill-climbing search: 8-queens problem


---



- $h$  = number of pairs of queens that are attacking each other, either directly or indirectly
- $h = 17$  for the above state
- The figure also shows the values of all successors, top successors have  $h = 12$

# Hill-climbing search: 8-queens problem

---



- A local minimum with  $h = 1$  (obtained in 5 steps)

# Hill-climbing search

---

- The algorithm gets stuck for several reasons:
  - Local Maximum: it is a peak that is higher than each of its neighbours, but lower than the maximum overall
  - Ridges: cause a sequence of local maxima that make navigation difficult
  - Plateau (flat): can lead to a local maximum where there is no ascendant exit or a terrace to advance
- In the 8-queens, it gets stuck in 86% and solve 14% cases
- If we allow lateral movements with the hope that we find a terrace (limiting them if reach a local maximum, e.g.100):→ 94% success
- Variants:
  - Stochastically (randomly chooses upward movements)
  - Random restart (the initial states are generated randomly)



# Informed search strategies

---

- ☐ Best-first search
- ☐ Greedy best-first search
- ☐ A\* search
- ☐ Heuristics
- ☐ Local search algorithms
  - Hill-climbing search
  - **Simulated annealing search**
  - Local beam search
  - Genetic algorithms

# Simulated annealing search

---

- ❑ Simulated annealing is the process of tempering or hardening metals by heating and then cooling them gradually
- ❑ Idea: escape local maxima by allowing some "bad" moves but **gradually decrease** their frequency
- ❑ It combines hill-climbing with random generation successor
- ❑ One can prove: If  $T$  decreases slowly enough, then simulated annealing search will find a global optimum with probability approaching 1
- ❑ Widely used in VLSI layout, airline scheduling, etc

# Informed search strategies

---

- ☐ Best-first search
- ☐ Greedy best-first search
- ☐ A\* search
- ☐ Heuristics
- ☐ Local search algorithms
  - Hill-climbing search
  - Simulated annealing search
  - **Local beam search**
  - Genetic algorithms

# Local beam search

---

- Idea: Keep track of  $k$  states rather than just one
  - Start with  $k$  randomly generated states
  - At each iteration, all the successors of all  $k$  states are generated
  - If any one is a goal state, stop; else select the  $k$  best successors from the complete list and repeat
  - Alternatively stochastic LBS randomly choose  $k$  successors, with the probability of choosing a successor as an increasing function of its value

# Informed search strategies

---


- ☐ Best-first search
- ☐ Greedy best-first search
- ☐ A\* search
- ☐ Heuristics
- ☐ Local search algorithms
  - Hill-climbing search
  - Simulated annealing search
  - Local beam search
  - **Genetic algorithms**

# Genetic algorithms

---

- A successor state is generated by combining two parent states
- Start with  $k$  randomly generated states (**population**)
- A state is represented as a string over a finite alphabet (often a string of 0s and 1s)
- Evaluation function (**fitness function**). Higher values for better states
- Produce the next generation of states by selection, crossover, and mutation


# Genetic algorithms



- The initial population is represented by chains of 8 digits representing states of 8 queens

# Genetic algorithms


---



- It corresponds to the 2 parents of figure (c)  
327 52411 (5 attacked queens)  
247 48552 (4 attacked queens)



# Genetic algorithms



- Fitness function: number of NON-attacking pairs of queens (min = 0, max =  $8 \times 7/2 = 28$ )
- Right: 247 48552 (4 attacked queens) =  $28 - 4 = 24$   
 $24/(24+23+20+11) = 31\%$
- Left: 327 52411 (5 attacked queens) =  $28 - 5 = 23$   
 $23/(24+23+20+11) = 29\%$

# Genetic algorithms



probabilities

- ❑ Crossing points are chosen and d) the offspring are created crossing the chain parental crosspoint
- ❑ In e) each position is subject to random mutation with a small probability

# Outline

---

- ☐ Introduction
- ☐ Problem-solving agents
- ☐ Problem formulation
- ☐ Problem types
- ☐ Example problems
- ☐ Basic search algorithms
- ☐ **Conclusions**

# Conclusions

---

- Problem formulation usually requires abstracting away real-world details to define a state space that can feasibly be explored
- Variety of uninformed and informed search strategies